

Exercise 02a: Map Reduce applications for Word Counting

Previous exercise described how to save input file in to HDFS. This exercise train students to do MapReduce process using word counting application.

Prerequisites

Ensure that Hadoop is installed, configured and is running. More details:

Single Node Setup for first-time users.

Cluster Setup for large, distributed clusters.

MapReduce Overview

Hadoop MapReduce is a software framework for easily writing applications which process vast amounts of data (multi-terabyte data-sets) in-parallel on large clusters (thousands of nodes) of commodity hardware in a reliable, fault-tolerant manner.

A MapReduce *job* usually splits the input data-set into independent chunks which are processed by the *map tasks* in a completely parallel manner. The framework sorts the outputs of the maps, which are then input to the *reduce tasks*. Typically both the input and the output of the job are stored in a file-system. The framework takes care of scheduling tasks, monitoring them and re-executes the failed tasks.

Typically the compute nodes and the storage nodes are the same, that is, the MapReduce framework and the Hadoop Distributed File System are running on the same set of nodes. This configuration allows the framework to effectively schedule tasks on the nodes where data is already present, resulting in very high aggregate bandwidth across the cluster.

The MapReduce framework consists of a single master `ResourceManager`, one worker `NodeManager` per cluster-node, and `MRAppMaster` per application.

Minimally, applications specify the input/output locations and supply *map* and *reduce* functions via implementations of appropriate interfaces and/or abstract-classes. These, and other job parameters, comprise the *job configuration*.

The Hadoop *job client* then submits the job (jar/executable etc.) and configuration to the `ResourceManager` which then assumes the responsibility of distributing the software/configuration to the workers, scheduling tasks and monitoring them, providing status and diagnostic information to the job-client.

Inputs and Outputs

The MapReduce framework operates exclusively on `<key, value>` pairs, that is, the framework views the input to the job as a set of `<key, value>` pairs and produces a set of `<key, value>` pairs as the output of the job, conceivably of different types.

The `key` and `value` classes have to be serializable by the framework and hence need to implement the `Writable` interface. Additionally, the key classes have to implement the `WritableComparable` interface to facilitate sorting by the framework.

Input and Output types of a MapReduce job:

(input) `<k1, v1>` -> **map** -> `<k2, v2>` -> **combine** -> `<k2, v2>` -> **reduce** -> `<k3, v3>` (output)

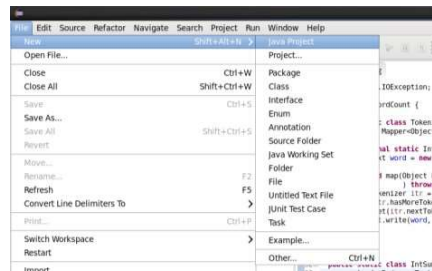
Step 1

Compile `WordCount.java` and create a jar:



- (i) Open Eclipse in Cloudera




(ii) Create 'WordCount' java project

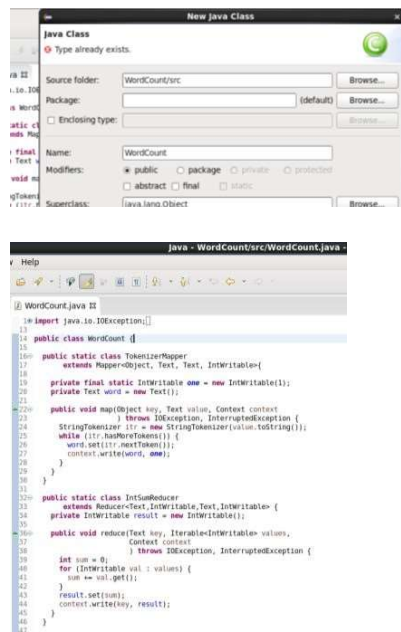


Import following Jar files

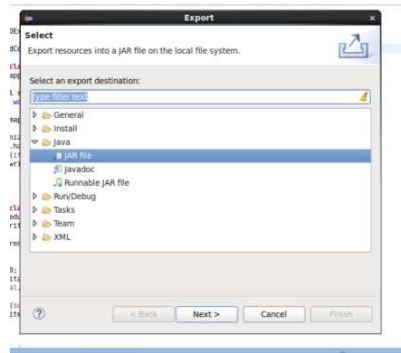


(iii) Create 'WordCount.java' in src folder





(iv) Create WordCount.jar file



Step 2

Create following folders in HDFS:

- /input - input directory in HDFS
- /output - output directory in HDFS

```

File Edit View Search Terminal Help
[cloudera@quickstart ~]$ ls
aaa.txt          eclipse          parcels          TempFile
cloudera-manager enterprise-deployment.json Pictures         Templates
cm_api.py        express-deployment.json Public          Videos
Desktop          kerberos       sujith.txt      WordCount.jar
Documents        lib            te              workspace
Downloads        Music          temp
[cloudera@quickstart ~]$ hdfs dfs -mkdir /in00

```

Step 3

Create and copy sample text-files into input folder:

```
[cloudera@quickstart ~]$ hdfs dfs -ls /in00/
```

Found 1 items

```
-rw-r--r-- 1 cloudera supergroup      158 2021-08-15 04:32 /in00/WCFile.txt
```

```
[cloudera@quickstart ~]$ cat > kausalya.txt
```

```
live your life
```

```
^C
```

```
[cloudera@quickstart ~]$ hdfs dfs -put kausalya.txt /in00/
```

Step 4

Run the MapReduce application:

```
hadoop jar /home/cloudera/WordCount.jar WordCount /in00/WCFile.txt /out00
```

Show MapReduce Framework

```
[cloudera@quickstart ~]$ hadoop jar /home/cloudera/WordCount.jar WordCount /in00/kausalya.txt /out19
22/08/03 08:00:00 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
22/08/03 08:00:00 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
22/08/03 08:00:00 INFO input.FileInputFormat: Total input paths to process : 1
22/08/03 08:00:00 INFO mapreduce.JobSubmitter: number of splits:1
22/08/03 08:00:00 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1659506889358_0008
22/08/03 08:00:00 INFO impl.YarnClientImpl: Submitted application application_1659506889358_0008
22/08/03 08:00:00 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1659506889358_0008/
22/08/03 08:00:00 INFO mapreduce.Job: Running job: job_1659506889358_0008
22/08/03 08:00:07 INFO mapreduce.Job: Job job_1659506889358_0008 running in uber mode : false
22/08/03 08:00:07 INFO mapreduce.Job:  map 0% reduce 0%
22/08/03 08:00:11 INFO mapreduce.Job:  map 100% reduce 0%
22/08/03 08:00:17 INFO mapreduce.Job:  map 100% reduce 100%
22/08/03 08:00:18 INFO mapreduce.Job: Job job_1659506889358_0008 completed successfully
22/08/03 08:00:18 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=39
    FILE: Number of bytes written=286759
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=130
    HDFS: Number of bytes written=21
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
```

Step 5

Output:

```
[cloudera@quickstart ~]$ hdfs dfs -ls /out00/
```

Found 2 items

```
-rw-r--r-- 1 cloudera supergroup      0 2021-08-15 04:41 /out00/_SUCCESS
```

```
-rw-r--r-- 1 cloudera supergroup 113 2021-08-15 04:41 /out00/part-r-00000
```

```
[cloudera@quickstart ~]$ hdfs dfs -cat /out00/part-r-00000
```

```
      Reduce shuffle bytes=39
      Reduce input records=3
      Reduce output records=3
      Spilled Records=6
      Shuffled Maps =1
      Failed Shuffles=0
      Merged Map outputs=1
      GC time elapsed (ms)=49
      CPU time spent (ms)=1010
      Physical memory (bytes) snapshot=499277824
      Virtual memory (bytes) snapshot=3137150976
      Total committed heap usage (bytes)=360710144
Shuffle Errors
      BAD_ID=0
      CONNECTION=0
      IO_ERROR=0
      WRONG_LENGTH=0
      WRONG_MAP=0
      WRONG_REDUCE=0
File Input Format Counters
      Bytes Read=16
File Output Format Counters
      Bytes Written=21
[cloudera@quickstart ~]$ hdfs dfs -ls /out19/
Found 2 items
-rw-r--r-- 1 cloudera supergroup      0 2022-08-03 08:00 /out19/_SUCCESS
-rw-r--r-- 1 cloudera supergroup    21 2022-08-03 08:00 /out19/part-r-00000
[cloudera@quickstart ~]$ hdfs dfs -cat /out19/part-r-00000
life    1
live    1
your    1
[cloudera@quickstart ~]$ █
```