

NAME:R.RAHINI ROLL NO:215229128

Lab8. Audio corpus creation and binary classification using DNN

1. Create a Dataset

Dataset is created

2. Read the Audio

In [2]:

```
import warnings
warnings.filterwarnings('ignore')

import os

import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
%matplotlib inline

import glob

from IPython.display import Audio
import IPython.display as ipd

import librosa #pip install Librosa
import librosa.display

from tqdm import tqdm

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten
from tensorflow.keras.optimizers import Adam, RMSprop
from tensorflow.keras.utils import to_categorical

from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import *
```

In [3]:


```
paths = '1fly.wav'  
Audio(paths)
```

Out[3]:

0:00 / 0:00

In [4]:

```
data, sample_rate = librosa.load(paths)  
plt.figure(figsize=(14,5))  
  
librosa.display.waveshow(data, sr=sample_rate)  
plt.show()
```



In [5]:

```
sample_rate
```

Out[5]:

22050

In [6]:

```
data
```

Out[6]:

```
array([0.          , 0.          , 0.          , ..., 0.01374257, 0.01699779,  
       0.          ], dtype=float32)
```

In [7]:

```
stftt = librosa.feature.chroma_stft(y=data, sr=sample_rate)  
stftt.shape
```

Out[7]:

(12, 105)

In [8]:

stftt

Out[8]:

```
array([[0.          , 0.          , 0.          , ..., 0.802695   , 1.          ,
        1.          ],
       [0.          , 0.          , 0.          , ..., 0.7765452   , 0.9367392   ,
        0.87225854],
       [0.          , 0.          , 0.          , ..., 0.64823526 , 0.67778885 ,
        0.70276004],
       ...,
       [0.          , 0.          , 0.          , ..., 0.62063694 , 0.42907584 ,
        0.4647396  ],
       [0.          , 0.          , 0.          , ..., 0.5343908   , 0.31281394 ,
        0.3764422  ],
       [0.          , 0.          , 0.          , ..., 0.7833488   , 0.6150779   ,
        0.60883486]], dtype=float32)
```

In [10]:

```
def features_extractor(file):
    audio, sample_rate = librosa.load(file_name)
    stftt_features = librosa.feature.chroma_stft(y=audio, sr=sample_rate)
    stftt_scaled_features = np.mean(stftt_features.T, axis=0)
    return stftt_scaled_features
```

In [11]:

```
extracted_features=[]

for index_num, row in tqdm(df.iterrows()):
    file_name = row[0]
    final_class_labels = row[1]
    data = features_extractor(file_name)
    extracted_features.append([data, final_class_labels])
```

20it [00:05, 3.52it/s]

In [12]:

```
extracted_features_df = pd.DataFrame(extracted_features, columns=['feature', 'class'])
extracted_features_df.head()
```

Out[12]:

	feature	class
0	[0.36663052, 0.38500747, 0.46057516, 0.5782328...	0
1	[0.33804783, 0.36386037, 0.39901736, 0.4628456...	0
2	[0.525747, 0.41264907, 0.22864214, 0.3536954, ...	0
3	[0.34989068, 0.30691823, 0.37625653, 0.2727835...	0
4	[0.33986202, 0.34667534, 0.28150865, 0.2368082...	0

3. Split the dataset

In [13]:

```
X = np.array(extracted_features_df['feature'].tolist())  
y = np.array(extracted_features_df['class'])
```

In [14]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=0)
```

In [15]:

```
print(X_train.shape)  
print(y_train.shape)  
print(X_test.shape)  
print(y_test.shape)
```

(15, 12)

(15,)

(5, 12)

(5,)

4. Train a Neural Network Model

In [16]:

```
batch_size=132  
num_labels = y.shape[0]
```

In [18]:

```

model=Sequential()
model.add(Dense(128, activation='tanh', input_shape=(12,)))
model.add(Dense(64, activation='tanh'))
model.add(Dense(32, activation='tanh'))
model.add(Dense(16, activation='tanh'))
model.add(Dense(8, activation='tanh'))
model.add(Dense(1, activation='sigmoid'))
model.summary()
model.compile(loss='mean_squared_error',metrics=['accuracy'],optimizer='adam')
history=model.fit(X_train, y_train, batch_size=batch_size, epochs=50 , verbose=2,validat

```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	1664
dense_1 (Dense)	(None, 64)	8256
dense_2 (Dense)	(None, 32)	2080
dense_3 (Dense)	(None, 16)	528
dense_4 (Dense)	(None, 8)	136
dense_5 (Dense)	(None, 1)	9
Total params: 12,673		
Trainable params: 12,673		

In [20]:

```



score=model.evaluate(X_test,y_test,verbose=0)
print("Loss :", score[0])
print("Accuracy :",score[1])

```

Loss : 0.19610339403152466
Accuracy : 0.800000011920929

In [21]:

```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['Train', 'Validation'])
plt.show()
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['Train', 'Validation'])
plt.show()
```



In [22]:

```
y_pred=model.predict(X_test)
y_pred=(y_pred>0.5)*1
y_pred
```

1/1 [=====] - 0s 234ms/step

Out[22]:

```
array([[1],
       [0],
       [0],
       [0],
       [1]])
```

In [23]:

```
y_pred = model.predict(X_test).round()
y_pred
```

1/1 [=====] - 0s 56ms/step

Out[23]:

```
array([[1.],
       [0.],
       [0.],
       [0.],
       [1.]], dtype=float32)
```

In [24]:

```
print(" Accuracy ",accuracy_score(y_test,y_pred))
print(" Precision ",precision_score(y_test,y_pred))
print(" Recall   ",recall_score(y_test,y_pred))
print(" AUC     ",roc_auc_score(y_test,y_pred))
```

```
Accuracy  0.8
Precision  1.0
Recall    0.6666666666666666
AUC       0.8333333333333333
```

5. Run different Neural Network models

In [25]:

```

def c_model(node):
    model=Sequential()
    model.add(Dense(128, activation='tanh', input_shape=(12,)))
    model.add(Dense(node, activation='tanh'))
    model.add(Dense(node, activation='tanh'))
    model.add(Dense(node, activation='tanh'))
    model.add(Dense(8, activation='tanh'))
    model.add(Dense(1, activation='sigmoid'))
    model.summary()
    model.compile(loss='mean_squared_error',metrics=['accuracy'],optimizer='adam')
    history=model.fit(X_train, y_train, batch_size=batch_size, epochs=50 , verbose=2,va
    score=model.evaluate(X_test,y_test,verbose=0)
    print("loss ", score[0])
    print("accuracy ",score[1])
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
    plt.title('Model Accuracy')
    plt.ylabel('accuracy')
    plt.xlabel('epoch')
    plt.legend(['Train', 'Validation'])
    plt.show()
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
    plt.title('Model loss')
    plt.ylabel('loss')
    plt.xlabel('epoch')
    plt.legend(['Train', 'Validation'])
    plt.show()

```

In [26]:

c_model(8)

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_6 (Dense)	(None, 128)	1664
dense_7 (Dense)	(None, 8)	1032
dense_8 (Dense)	(None, 8)	72
dense_9 (Dense)	(None, 8)	72
dense_10 (Dense)	(None, 8)	72
dense_11 (Dense)	(None, 1)	9

```

=====
Total params: 2,921
Trainable params: 2,921
Non-trainable params: 0

```


In [27]:

```
c_model(16)
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
dense_12 (Dense)	(None, 128)	1664
dense_13 (Dense)	(None, 16)	2064
dense_14 (Dense)	(None, 16)	272
dense_15 (Dense)	(None, 16)	272
dense_16 (Dense)	(None, 8)	136
dense_17 (Dense)	(None, 1)	9
=====		
Total params: 4,417		
Trainable params: 4,417		
Non-trainable params: 0		

In [28]:

```
c_model(32)
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
dense_18 (Dense)	(None, 128)	1664
dense_19 (Dense)	(None, 32)	4128
dense_20 (Dense)	(None, 32)	1056
dense_21 (Dense)	(None, 32)	1056
dense_22 (Dense)	(None, 8)	264
dense_23 (Dense)	(None, 1)	9
=====		
Total params: 8,177		
Trainable params: 8,177		
Non-trainable params: 0		

In [29]:

c_model(64)

Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_24 (Dense)	(None, 128)	1664
dense_25 (Dense)	(None, 64)	8256
dense_26 (Dense)	(None, 64)	4160
dense_27 (Dense)	(None, 64)	4160
dense_28 (Dense)	(None, 8)	520
dense_29 (Dense)	(None, 1)	9

=====
Total params: 18,769
Trainable params: 18,769
Non-trainable params: 0

In [33]:

```
def c_layer(n):
    model=Sequential()
    model.add(Dense(128, activation='tanh', input_shape=(12,)))
    for i in range(0,n):
        model.add(Dense(32, activation='tanh'))

    model.add(Dense(1, activation='sigmoid'))
    model.summary()
    model.compile(loss='mean_squared_error',metrics=['accuracy'],optimizer='adam')
    history=model.fit(X_train, y_train, batch_size=batch_size, epochs=50 , verbose=2,validation_data=(X_test,y_test),
    score=model.evaluate(X_test,y_test,verbose=0)
    print("loss ", score[0])
    print("accuracy ",score[1])
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
    plt.title('Model Accuracy')
    plt.ylabel('accuracy')
    plt.xlabel('epoch')
    plt.legend(['Train', 'Validation'])
    plt.show()
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
    plt.title('Model loss')
    plt.ylabel('loss')
    plt.xlabel('epoch')
    plt.legend(['Train', 'Validation'])
    plt.show()
```

In [34]:

```
c_layer(2)
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
dense_30 (Dense)	(None, 128)	1664
dense_31 (Dense)	(None, 32)	4128
dense_32 (Dense)	(None, 32)	1056
dense_33 (Dense)	(None, 1)	33

=====
Total params: 6,881
Trainable params: 6,881
Non-trainable params: 0

Epoch 1/50
1/1 - 1s - loss: 0.2464 - accuracy: 0.5833 - val_loss: 0.2624 - val_accu

In [35]:

```
c_layer(3)
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
dense_34 (Dense)	(None, 128)	1664
dense_35 (Dense)	(None, 32)	4128
dense_36 (Dense)	(None, 32)	1056
dense_37 (Dense)	(None, 32)	1056
dense_38 (Dense)	(None, 1)	33

=====
Total params: 7,937
Trainable params: 7,937
Non-trainable params: 0

Epoch 1/50

In [36]:

```
c_layer(5)
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
=====		
dense_39 (Dense)	(None, 128)	1664
dense_40 (Dense)	(None, 32)	4128
dense_41 (Dense)	(None, 32)	1056
dense_42 (Dense)	(None, 32)	1056
dense_43 (Dense)	(None, 32)	1056
dense_44 (Dense)	(None, 32)	1056
dense_45 (Dense)	(None, 1)	33
=====		
Total params: 10,352		